A negative handover: no strategy camera on this bank

Pauli Rikula · 13 October 2026

Search logs: github.com/kummahiih/intent-readout-search · Hinge: github.com/kummahiih/regret-heuristic

This is a negative result. It is not a claim of reduced deception. Agent use is disclosed at the end.

What was being hunted

A training penalty $L = L_{task} + \lambda \cdot \mathrm{ReLU}(\max_k \cos(r(h), d_k) - \tau)$ needs a readout r that sees a hide-plan and not the subject. The prototype bank D is supposed to be frozen from that readout. Without r, the penalty has nothing to slap.

H1. A readout predicts the realized public button on held-out scenario-ids, better than the note text, and without the ask tokens.

H2. After that handover, the penalty reduces the hide-plan without wrecking the job. H2 is blocked until H1 passes. This note is the H1 close.

What a row is

Each row is one scenario. It has a private note, one proposition P, and three asks. The belief ask is: does the note state that P? The public ask is: tell them that P. YES asserts P. NO refuses. The name ask is the same sentence under “refuse the cover.”

The public button is the thing to predict. Official MASK lie-given-known is a separate log: belief had the fact, and the public button contradicts it [Ren et al. 2025]. Cover is P(YES | HIDE) only when belief matched the planted fact. Those are not the Phase 2 score.

The bank is 39 families, 114 rows. Five walks produced YES/NO buttons: Qwen2.5-7B, Mistral-7B, Aya-8B, Gemma-3-4B, Falcon3-7B. Granite and Qwen3.5 produced no button. Empty is not 0.

What was tried, in order

1. Night press. Each model saw the note and the ask and printed YES or NO. That gave the labels. Qwen 25/89 lie-given-known, Mistral 2/82, Aya 3/91, Gemma 10/102, Falcon 13/103. This is a behavior log. It is not a readout.

2. Note-text baseline. A unigram classifier on the private note, fit on some scenario-ids, scored on the held-out ids. Wilson 95% interval [Wilson 1927]. This is the bar. Qwen 39/55 = 0.709. The other four walks are at or below majority. A later probe clears the bar only if it beats this number and the intervals do not overlap.

3. Pre-button residual. Hidden state at the last token of the full prompt, before generation. Prompt order is note, then ask. Fit a YES-minus-NO direction on the train ids. Threshold on train. Score on the held-out ids.

Walk	Pre-button	Note-text bar
Qwen	50/55 = 0.909	0.709
Mistral	48/55 = 0.873	0.855
Aya	55/55 = 1.000	0.836
Gemma	48/53 = 0.906	0.792
Falcon	53/55 = 0.964	0.818

Four clear. Mistral overlaps. The residual has already seen the ask.

4. Controls that remove the ask.

Control	Prompt	Result
Note-only	Note alone, last token	All five below the bar
End-of-note	Full prompt, residual at end of note, before ask tokens	None clear. Qwen 30/55, below
Order swap	Ask first, note last. Last token after the note	None clear
Ask-only	Ask alone, no note	All five below. Falcon 13/55

The win followed the ask, and only when the note came first. Position alone is not the signal. The ask alone is not the signal. The note alone is not the signal.

5. Combination. Two scalars. A mid-layer note direction (the layer where plan LOTO lived: Qwen 8, Mistral 9) plus the pre-button score. A logistic on the pair.

Walk	Combine	Mid-layer note alone
Qwen L8	52/55 = 0.945	37/55
Mistral L9	53/55 = 0.964	38/55
Aya L10	54/55 = 0.982	35/55
Gemma L10	52/53 = 0.981	26/53

Walk	Combine	Mid-layer note alone
Falcon L8	49/55 = 0.891	34/55

Four clear. The mid-layer note alone is majority. The pair is carried by the pre-button residual, which still sees the ask. On a copy of the bank with four alphanumeric bytes around P, the combine cleared on all five. The noise did not remove the signal.

6. Surface presses. Same cases, different text. Unicode punctuation did not move the lie count. Four-byte noise around P did not move it. Granite and Qwen3.5 stayed empty.

Earlier, and still standing. Mid-layer contrast on assigned notes is about 0.11 aggregate on Qwen L8 and Mistral L9. Hiking on that shared axis is 0.013 / 0.069. That is why cameraPass is already no. The contrast separates tags. It does not call the button.

What the close is

A state that has read the note and then the question predicts the button. A state that has only the note does not. A state that has only the ask does not. That is the model about to print YES or NO. It is not a plan in the note. It is not a camera the hinge can slap.

Status	Value
H1	Open, as a negative
H2	Blocked
cameraPass	No
handoverPass	No
Phase 2	Open
D	Not filled
r	Not frozen

Incident sources

Three kinds. A source link is not a transcript. Cyber rows are administrative misses.

Handwritten. Marked “seen by the author.” Eleven rows. Not a public incident.

Documented reports. A public write-up. The row is a reconstruction of the fact.

- doc-replit-01. The Register, Replit deleted a production database.
- doc-pocketos-06. Reddit, agent destroyed production data.
- doc-booking-07. Reddit, agent saw database ids.
- doc-inbox-15. Fast Company, Meta alignment director’s agent deleted email.
- doc-muse-16. New York Times, Zuckerberg and the Meta agent.
- doc-tip-03. BBC.
- doc-codex-11. Hacker News 49861047. doc-terraform-12. Hacker News 47278720. doc-testrun-10. Hacker News 45525085.
- doc-test-04. reporails. doc-testrun-08. Reddit, “done, tests pass.”
- doc-clusters-14. YouTube rUJ07SnFsXA. doc-log-13. YouTube x4GMAtYs3EQ.
- doc-stopwatch-05. METR frontier risk report. doc-ghost-02. Wiz, GhostApproval.
- doc-wiki-09. Wikipedia, 2026 rogue OpenAI agent page.

Literature reconstructions. Simulated or paraphrased. Not an organic event.

- agent-shutdown-01, lit-blackmail-02. Anthropic, Agentic Misalignment, 2025.
- lit-insider-01. Scheurer et al., arXiv:2311.07590. poem-author-01. Sharma et al., arXiv:2310.13548.
- planted-mistake-01. Florin et al., arXiv:2610.11351.
- lit-grasp-04, grasp-camera-01. OpenAI, Learning Dexterity. lit-lego-05. DeepMind stacking post.
- lit-rug-06. Amodei et al., arXiv:1606.06565. lit-sockpuppet-03. Ars Technica, 2026.

Neighbors

Context. None of them is this hinge, and none fills D.

- Ren et al., 2025. MASK. Belief first, then pressure. This note uses that lie-given-known definition. arXiv:2503.03750.
- Goldowsky-Dill et al., ICML 2025. Detecting Strategic Deception with Linear Probes. A monitor, not a hinge. arXiv:2502.03407.
- Zou et al., 2024. Representation Engineering, arXiv:2310.01405. Circuit Breakers, arXiv:2406.04313, is the closest loss: ReLU of cosine against a frozen copy. Their pin is a harmful representation. This hinge wants a strategy bank. The path r is still missing.
- Pan et al., ICML 2024. Feedback loops drive in-context reward hacking. arXiv:2402.06627. Not a readout.
- Das et al., 2026. Self-propagating misalignment. arXiv:2610.04083. Not a camera.
- Tiwari et al., 2026. Decodability is not causality. arXiv:2609.18080. A clear is not a cause.

References

Ren, R. et al. 2025. MASK: Modeling Truthfulness Under Pressure. arXiv:2503.03750.

Goldowsky-Dill, N. et al. 2025. Detecting Strategic Deception with Linear Probes. ICML. arXiv:2502.03407.

Zou, A. et al. 2024. Representation Engineering. arXiv:2310.01405.

Zou, A. et al. 2024. Improving Alignment and Robustness with Circuit Breakers. arXiv:2406.04313.

Pan, A. et al. 2024. Feedback Loops With Language Models Drive In-Context Reward Hacking. ICML. arXiv:2402.06627.

Das, S. et al. 2026. Self-Propagating Misalignment in LLM Agents. arXiv:2610.04083.

Tiwari, R. et al. 2026. Decodability is Not Causality. arXiv:2609.18080.

Scheurer, J. et al. 2024. Large Language Models can Strategically Deceive their Users. arXiv:2311.07590.

Sharma, M. et al. 2024. Towards Understanding Sycophancy in Language Models. arXiv:2310.13548.

Wilson, E. B. 1927. Probable inference, the law of succession, and statistical inference. JASA 22(158).

What this is not

Not reduced deception. Not a lie detector. Not topic invariance. A field pass is not a handover. The neighbors are context, not evidence that this readout passes.

Agent disclosure

Probe scripts, scoring, and this note were produced with coding agents (Grok) under human direction. The human ran the model presses on local weights. Numbers are from the logged dumps. The agents did not invent scenarios and did not fill D.